

BUILDING A MODULAR MENTAL HEALTH RAG CHATBOT

Language Detection, Intent & Emotion Classification,
RAG, and Architecture | For Educational Audiences



TABLE OF CONTENTS

Roadmap to Building a Modular Mental Health RAG Chatbot

1



**Problem
Statement**

2



**Language
Detection**

3



**Intent
Classification**

4



**Emotion
Classification**

5



RAG

6



**Full
Architecture**



WHY A MENTAL HEALTH CHATBOT?

Mental health support is more critical than ever, yet access remains limited by geography, cost, and stigma. A mental health chatbot offers 24/7 availability, scalable support for millions of users, and a private, judgment-free space to seek help. By combining AI-driven language understanding with empathetic design, these systems bridge the gap between need and care — making mental wellness resources accessible to everyone, anywhere.



INTRODUCTION TO RAG CHATBOTS



RAG = Retrieval-Augmented Generation: combines retrieval of relevant knowledge with generative AI to produce accurate, context-aware responses.



Enhances chatbot responses with up-to-date, factual information — improving trust and relevance in sensitive domains like mental health support.



MODULAR ARCHITECTURE OVERVIEW

Language Detection & Intent Classification: Identifies the user's language for multilingual support, then determines the user's goal — whether seeking advice, resources, or urgent help — enabling personalized and relevant responses.

Emotion Classification & RAG: Detects the user's emotional state (e.g. anxiety, sadness, calm) to tailor empathetic replies, then uses Retrieval-Augmented Generation to ground responses in trusted mental health knowledge sources.

Full System Architecture & Modularity Benefits: All modules connect in a clear data flow pipeline. The modular design ensures scalability, easier maintenance, faster experimentation, and a clean separation of concerns for developers and educators.

MODULE 1: LANGUAGE DETECTION

Language Detection identifies the language of a user's input before any further processing.

- Supports multilingual users for personalized interaction
- Techniques: rule-based heuristics, ML classification, transformer models (e.g. XLM-R)
- Libraries: fastText, LangID for fast and accurate detection

Example: A user types "Necesito ayuda" — the module detects Spanish and routes the message to a Spanish-language response pipeline, ensuring culturally appropriate and accurate mental health support.



MODULE 2 – INTENT CLASSIFICATION

Intent classification identifies what the user wants from the conversation, enabling the chatbot to respond appropriately.

- Role: Determines user goals — seeking advice, requesting resources, or signaling urgent help
- Approaches: Supervised machine learning, neural networks, BERT-based classifiers
- Example intents:
 - "Need coping strategies"
 - "Request emergency contacts"
 - "Chat casually"

Accurate intent detection is critical for routing users to the right response and ensuring safety in sensitive mental health contexts.



MODULE 3 – EMOTION CLASSIFICATION

Emotion Classification detects the emotional state of the user to enable empathetic, appropriate responses.

- Purpose: Tailor chatbot responses based on detected emotions such as sadness, anxiety, or calm
- Methods: Sentiment analysis, fine-tuned emotion detection models (e.g. RoBERTa)
- Detected states: Sadness, Anxiety, Anger, Calm, Neutral
- Example: A message classified as high anxiety triggers a calming, supportive response with breathing exercises
- Ensures the chatbot responds with emotional intelligence and avoids harmful or dismissive replies



MODULE 4: RAG



Retrieval-Augmented Generation (RAG) combines a retrieval engine with a generative AI model. When a user sends a message, relevant documents are fetched from trusted mental health sources, then passed to the generative model to craft an accurate, context-aware, and empathetic response — grounding answers in real knowledge.



Traditional chatbot: relies solely on pre-trained patterns, may produce generic or outdated answers. RAG-enabled chatbot: retrieves up-to-date, domain-specific documents at query time, producing more reliable, informative, and contextually sensitive responses — critical in mental health support.




RAG EXAMPLE IN ACTION

User submits a mental health question through the chatbot interface. The query is preprocessed through language detection, intent classification, and emotion classification modules to understand context, language, and emotional state before retrieval begins.

Relevant documents are retrieved from a trusted mental health knowledge base using dense vector search. The retriever identifies the most semantically similar passages to the user's query, ensuring responses are grounded in accurate, up-to-date information.

A generative model crafts a natural, empathetic response by combining the retrieved documents with contextual understanding. The output ensures factual accuracy, emotional sensitivity, and a supportive tone tailored to the user's detected emotional state.

FULL ARCHITECTURE DIAGRAM



End-to-end data flow: User input passes through Language Detection, then Intent & Emotion Classification modules, feeding context-rich signals into the RAG retrieval and generation pipeline before delivering an empathetic, informed response output.

Modular separation ensures each component — language detection, intent classification, emotion classification, and RAG — operates independently yet interacts seamlessly, enabling scalability, easier maintenance, and clear boundaries for upgrades or debugging.




TECHNICAL COMPONENTS

Language Detection & Intent Classification: fastText and LangID for multilingual input detection. BERT-based classifiers for intent recognition — identifying user goals such as seeking advice, requesting resources, or urgent help.


Emotion Classification & RAG: RoBERTa fine-tuned on emotion datasets (sadness, anxiety, calm) for empathetic response tailoring. Dense retrievers paired with GPT-style generative models for Retrieval-Augmented Generation.

System Architecture: Microservices-based design with REST APIs connecting each module. Enables independent scaling, easier maintenance, and clear separation of concerns across language, intent, emotion, and RAG layers.

BENEFITS OF MODULAR DESIGN



Flexibility to upgrade or replace individual modules independently without affecting the entire system. Easier debugging and maintenance — isolate issues within a single module for faster resolution and cleaner codebases.



Faster experimentation and customization — swap models or techniques per module with minimal friction. Clear separation of concerns makes the system more understandable for both developers and educators building on top of it.

CHALLENGES & CONSIDERATIONS



Handling ambiguous or mixed-language inputs can reduce accuracy across modules. Ensuring emotional sensitivity is critical – responses must avoid causing harm, especially for users in distress or crisis situations.



Maintaining an up-to-date and accurate retrieval corpus requires ongoing curation of trusted mental health sources. Privacy and data security must be prioritized to protect sensitive user conversations and comply with regulations.

FUTURE DIRECTIONS

Looking ahead, the modular chatbot can evolve in exciting ways. Multimodal inputs such as voice and images will broaden accessibility. Deeper personalization through user profiles will tailor responses more meaningfully. Expanded language support and cultural adaptation will serve diverse global populations. Finally, seamless integration with human expert escalation ensures safety-critical cases receive professional attention.



SUMMARY & KEY TAKEAWAYS



The modular mental health RAG chatbot architecture demonstrates how Language Detection, Intent Classification, Emotion Classification, and RAG can work together seamlessly. Each module plays a critical, well-defined role — enabling scalability, maintainability, and clarity across the full system.



RAG significantly improves response quality and reliability by grounding generative AI in trusted, up-to-date mental health knowledge. This educational exploration into building complex ML systems modularly offers a blueprint for responsible, empathetic, and effective AI in sensitive domains.

System Architecture

THANK YOU!

Questions? | Further resources available upon request

